

# Financial Simulation using R

Majeed Simaan<sup>1</sup>

<sup>1</sup>Lally School of Management at RPI

Feb 2, 2018

# Agenda

- Crash Course in R
  - Basics
  - Loops and the `apply` functions
  - Data Manipulation
  - Plotting
- Simulations
  - Distributions in R
  - Generating Random Numbers
  - Density Functions and Moments
- Option Pricing
  - Simulating Stock Prices
  - Pricing Options

# Agenda

- Crash Course in R
  - Basics
  - Loops and the `apply` functions
  - Data Manipulation
  - Plotting
- Simulations
  - Distributions in R
  - Generating Random Numbers
  - Density Functions and Moments
- Option Pricing
  - Simulating Stock Prices
  - Pricing Options
- If time allows...
  - Open Source Data (`quantmod` package)
  - Financial Time Series (`xts` package)

# Suggested Readings and Resources

## • Intro

- 1 Lally School January 2015 R Bootcamp (slides)
- 2 The R manuals (see link)
- 3 DataCamp
- 4 The Art of R Programming by Matloff (2011) (see link)

# Suggested Readings and Resources

- Intro

- 1 Lally School January 2015 R Bootcamp (slides)
- 2 The R manuals (see link)
- 3 DataCamp
- 4 The Art of R Programming by Matloff (2011) (see link)

- Advanced

- 1 The R Inferno by Burns (2011) (see link)
- 2 Advanced R by Wickham (2014) (see link)

# Suggested Readings and Resources

## • Intro

- 1 Lally School January 2015 R Bootcamp (slides)
- 2 The R manuals (see link)
- 3 DataCamp
- 4 The Art of R Programming by Matloff (2011) (see link)

## • Advanced

- 1 The R Inferno by Burns (2011) (see link)
- 2 Advanced R by Wickham (2014) (see link)

## • Applications

- 1 R for Data Science by Golemund and Wickham (2016) (see link)
- 2 Statistical Learning by James et al., 2014
- 3 Financial Modeling by Ang, 2015
- 4 Analysis of Financial Time Series by Tsay (2010) (see link)

# Introduction

# Types of Data

- R has five basic classes
  - 1 Logical
  - 2 Integer
  - 3 Numeric
  - 4 Character
  - 5 Complex
- The most basic object is a **vector**
  - Can only contain objects of the same class
  - The same applies to matrices
- A **list** can have objects of different types
- **In R, not everything is a matrix!**



# Basics

- Can use either `<-` or `=` to assign new objects

```
> x <- 1:5
```

```
> x
```

```
[1] 1 2 3 4 5
```

# Basics

- Can use either `<-` or `=` to assign new objects

```
> x <- 1:5  
> x
```

```
[1] 1 2 3 4 5
```

- Can also assign labels to objects

```
> names(x) <- letters[1:5] # assign names to x  
> x
```

```
a b c d e  
1 2 3 4 5
```

- Note that R has a number of built-in objects, e.g. `letters`

# Basics

- Can use either `<-` or `=` to assign new objects

```
> x <- 1:5  
> x  
  
[1] 1 2 3 4 5
```

- Can also assign labels to objects

```
> names(x) <- letters[1:5] # assign names to x  
> x  
  
a b c d e  
1 2 3 4 5
```

- Note that R has a number of built-in objects, e.g. `letters`

- Can pull data using indices or labels, e.g.

```
> x[1:2] # pulls the first two elements  
  
a b  
1 2  
  
> x["c"] # pulls the element labeled c  
  
c  
3
```

- Returns `NA` if label/index does not exist, e.g. `x["z"]`

- Note that  $x$  is a vector but not a matrix

```
> dim(x)
```

```
NULL
```

- If we define  $x$  as a matrix

```
> X <- as.matrix(x)
```

```
> dim(X)
```

```
[1] 5 1
```

- Note that `x` is a vector but not a matrix

```
> dim(x)
```

```
NULL
```

- If we define `x` as a matrix

```
> X <- as.matrix(x)
```

```
> dim(X)
```

```
[1] 5 1
```

- R uses the function `c` to combine vectors or lists

```
> x2 <- c(x,x,x,x,x)
```

```
> x2
```

```
a b c d e a b c d e a b c d e a b c d e a b c d e  
1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5
```

- Note that `x` is a vector but not a matrix

```
> dim(x)
```

```
NULL
```

- If we define `x` as a matrix

```
> X <- as.matrix(x)
```

```
> dim(X)
```

```
[1] 5 1
```

- R uses the function `c` to combine vectors or lists

```
> x2 <- c(x,x,x,x,x)
```

```
> x2
```

```
a b c d e a b c d e a b c d e a b c d e a b c d e
1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5
```

- You can also stack vectors in a matrix using the `matrix` function

```
> X2 <- matrix(x2,length(x))
```

```
> X2
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     1     1     1     1     1
[2,]     2     2     2     2     2
[3,]     3     3     3     3     3
[4,]     4     4     4     4     4
[5,]     5     5     5     5     5
```

- R has a repetitive nature

```
> x + 1:10
```

```
[1] 2 4 6 8 10 7 9 11 13 15
```

- whereas

```
> x + 1:9
```

```
[1] 2 4 6 8 10 7 9 11 13
```

Warning message:

In x + 1:9 :

longer object length is not a multiple of shorter object length

- R has a repetitive nature

```
> x + 1:10  
[1] 2 4 6 8 10 7 9 11 13 15
```

- whereas

```
> x + 1:9  
[1] 2 4 6 8 10 7 9 11 13  
Warning message:  
In x + 1:9 :  
  longer object length is not a multiple of shorter object length
```

- One benefit of repetitions is creating larger dimensions using smaller one

```
# create matrix using the vector x  
> X3 <- matrix(x,length(x),length(x))  
  
> identical(X3,X2) # returns TRUE if two objects are identical  
  
[1] TRUE
```



- Similar to vectors, rows and columns of matrices can be assigned labels

```
> rownames(X3) <- letters[1:nrow(X3)]  
> colnames(X3) <- LETTERS[1:ncol(X3)]  
> X3.subset <- X3[c("a", "c"), c("D", "E")]  
> X3.subset
```

```
  D E  
a 1 1  
c 3 3
```

- Note that . is not different than any other letter

- Similar to vectors, rows and columns of matrices can be assigned labels

```
> rownames(X3) <- letters[1:nrow(X3)]
> colnames(X3) <- LETTERS[1:ncol(X3)]
> X3.subset <- X3[c("a", "c"), c("D", "E")]
> X3.subset
```

```
  D E
a 1 1
c 3 3
```

- Note that `.` is not different than any other letter
- Alternatively, one can use indices to pull items from a matrix

```
> X3.subset2 <- X3[c(1,3), c(4,5)]
> all(X3.subset == X3.subset2)

[1] TRUE
```

- Also, to refer to a row/column, it suffices to call either dimension, e.g.

```
> X3[1,]

A B C D E
1 1 1 1 1

> X3[,1]

a b c d e
1 2 3 4 5
```

## Practice

- 1 Write a one line command that stacks all numbers from 1 to 100 into a  $10 \times 10$  squared matrix named `M`, such that

> `M`

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
[1,]	1	2	3	4	5	6	7	8	9	10
[2,]	11	12	13	14	15	16	17	18	19	20
[3,]	21	22	23	24	25	26	27	28	29	30
[4,]	31	32	33	34	35	36	37	38	39	40
[5,]	41	42	43	44	45	46	47	48	49	50
[6,]	51	52	53	54	55	56	57	58	59	60
[7,]	61	62	63	64	65	66	67	68	69	70
[8,]	71	72	73	74	75	76	77	78	79	80
[9,]	81	82	83	84	85	86	87	88	89	90
[10,]	91	92	93	94	95	96	97	98	99	100

**Hint:** note that by default

```
matrix(data = NA, nrow = 1, ncol = 1, byrow = FALSE, dimnames =
NULL)
```

- 2 Use common summary functions you already from MATLAB to summarize the `M` matrix, e.g. `min`, `max`, `mean`, etc...

**Hint:** note that common functions have the same name in MATLAB and R (see the R/MATLAB manual for a reference <https://cran.r-project.org/doc/contrib/Hiebeler-matlabR.pdf>)

# Functions

- Unlike MATLAB, R does not require functions to be stored in separate files
- Multiple functions can be defined in the main script

# Functions

- Unlike MATLAB, R does not require functions to be stored in separate files
- Multiple functions can be defined in the main script
- To define a function  $f$  that, for instance, returns the first and last elements of a given vector  $x$ , one can write

```
> f1 <- function(x) {  
+   x1 <- x[1]  
+   x2 <- x[length(x)]  
+   return(c(x1,x2)) # functions must return something  
+ }  
>  
> f(x)  
  
a e  
1 5
```
- Note that the braces  $\{ \}$  are needed to set the start and end of a function

# Functions

- Unlike MATLAB, R does not require functions to be stored in separate files
- Multiple functions can be defined in the main script
- To define a function  $f$  that, for instance, returns the first and last elements of a given vector  $x$ , one can write

```
> f1 <- function(x) {  
+   x1 <- x[1]  
+   x2 <- x[length(x)]  
+   return(c(x1,x2)) # functions must return something  
+ }  
>  
> f(x)  
  
a e  
1 5
```

- Note that the braces  $\{ \}$  are needed to set the start and end of a function
- Similar to the  $@(x)$  command in MATLAB, one can write

```
> f2 <- function(x) c(x[1],x[length(x)])  
> f2(x)  
  
a e  
1 5
```

- Functions in R can load objects defined outside the function, for instance

```
> f3 <- function(x) mean((x - mean(x)) > a*sd(x))  
> a <- 0.5  
> f3(x)
```

```
[1] 0.4
```

```
> a <- 1  
> f3(x)
```

```
[1] 0.2
```

- However, this also increases the chances to commit an error in the code

- Functions in R can load objects defined outside the function, for instance

```
> f3 <- function(x) mean((x - mean(x)) > a*sd(x))  
> a <- 0.5  
> f3(x)
```

```
[1] 0.4
```

```
> a <- 1  
> f3(x)
```

```
[1] 0.2
```

- However, this also increases the chances to commit an error in the code

### Practice

Write a function named `M.f` that takes two arguments: matrix `M` and constant `a`. The function should return the proportion of elements that are larger than `a`. For instance, the function should return

```
> M.f(M, -Inf)
```

```
[1] 1
```

```
> M.f(M, Inf)
```

```
[1] 0
```

where R defines `Inf` as an infinite number



# Loops

- Similar to MATLAB, R has `for` and `while` loops
- The `for` loop

```
> M.mean2 <- numeric() # define an empty numeric object
> for(i in 1:ncol(M)) {
+   mean.i <- mean(M[,i])
+   M.mean2 <- c(M.mean2, mean.i)
+ }
> M.mean2

[1] 46 47 48 49 50 51 52 53 54 55
```

# Loops

- Similar to MATLAB, R has for and while loops

- The for loop

```
> M.mean2 <- numeric() # define an empty numeric object
> for(i in 1:ncol(M)) {
+   mean.i <- mean(M[,i])
+   M.mean2 <- c(M.mean2, mean.i)
+ }
> M.mean2
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```

- The while loop

```
> i <- 1
> M.mean2 <- numeric() # define an empty numeric object
> while (i <= ncol(M)) {
+   mean.i <- mean(M[,i])
+   M.mean2 <- c(M.mean2, mean.i)
+   i <- i + 1
+ }
> M.mean2
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```

# The apply functions

- It is recommended to avoid loops when one can
- Fortunately, R provides a number of good alternatives

# The apply functions

- It is recommended to avoid loops when one can
- Fortunately, R provides a number of good alternatives
  - ① `apply`: applies a function over rows/columns of a given matrix, e.g.

```
> apply(M, 1, mean) # for rows
```

```
[1]  5.5 15.5 25.5 35.5 45.5 55.5 65.5 75.5 85.5 95.5
```

```
> apply(M, 2, mean) # for columns
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```

# The apply functions

- It is recommended to avoid loops when one can
- Fortunately, R provides a number of good alternatives
  - ① `apply`: applies a function over rows/columns of a given matrix, e.g.

```
> apply(M, 1, mean) # for rows
```

```
[1]  5.5 15.5 25.5 35.5 45.5 55.5 65.5 75.5 85.5 95.5
```

```
> apply(M, 2, mean) # for columns
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```

- ② `sapply`: applies a function over a given set, e.g.

```
> sapply(1:ncol(M), function(i) mean(M[,i]))
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```

# The apply functions

- It is recommended to avoid loops when one can
- Fortunately, R provides a number of good alternatives
  - 1** `apply`: applies a function over rows/columns of a given matrix, e.g.
 

```
> apply(M, 1, mean) # for rows
```

```
[1] 5.5 15.5 25.5 35.5 45.5 55.5 65.5 75.5 85.5 95.5
```

```
> apply(M, 2, mean) # for columns
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```
  - 2** `sapply`: applies a function over a given set, e.g.
 

```
> sapply(1:ncol(M), function(i) mean(M[,i]))
```

```
[1] 46 47 48 49 50 51 52 53 54 55
```
- Using either, one can define a temporary function to be applied:
 

```
> apply(M, 2, function(x) mean(((x - mean(x))/sd(x))^2))
```

```
[1] 0.9 0.9 0.9 0.9 0.9 0.9 0.9 0.9 0.9 0.9
```

## Practice

The `EuStockMarkets` object is a built-in dataset that contains data on 4 major European stock indices. You will need to

- ❶ store the data into an object named `ds`
- ❷ use a loop to compute the min and max of each index in the data
  - report the results in a  $2 \times 4$  matrix, where columns are named correspondingly
- ❸ repeat the above using the `apply` function
  - can you do it in one line?
- ❹ finally, how much time did you save by using the `apply` function?
  - **hint** use the `Sys.time()` command to capture the system's time

# Data Frames

- The problem with matrices is that all columns must have the same format

```
> i <- 1:26  
> l <- letters  
> L <- cbind(i,l) # combines columns  
> sapply(1:ncol(L),function(i) class(L[,i]) )  
  
           i           l  
"character" "character"
```

- A list object can combine objects of different formats



# Data Frames

- The problem with matrices is that all columns must have the same format

```
> i <- 1:26
> l <- letters
> L <- cbind(i,l) # combines columns
> sapply(1:ncol(L),function(i) class(L[,i]) )

      i      l
"character" "character"
```

- A list object can combine objects of different formats

- Nevertheless, data.frame objects provide a better solution

```
> L2 <- data.frame(i,l, stringsAsFactors = F) # avoid factors
> sapply(1:ncol(L2),function(i) class(L2[,i]) )

[1] "integer"    "character"

> summary(L2)

      i      l
Min.   : 1.00   Length:26
1st Qu.: 7.25   Class :character
Median :13.50   Mode  :character
Mean   :13.50
3rd Qu.:19.75
Max.   :26.00
```

- The summary is a “cheap” way to look at the data

- The `plyr` package (Wickham, 2011) is very useful to manipulate data frames (df)
- Let's take a look at the `iris` built-in data frame

```
> summary(iris,digits = 2)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.3	Min. :2.0	Min. :1.0	Min. :0.1	setosa :50
1st Qu.:5.1	1st Qu.:2.8	1st Qu.:1.6	1st Qu.:0.3	versicolor:50
Median :5.8	Median :3.0	Median :4.3	Median :1.3	virginica :50
Mean :5.8	Mean :3.1	Mean :3.8	Mean :1.2	
3rd Qu.:6.4	3rd Qu.:3.3	3rd Qu.:5.1	3rd Qu.:1.8	
Max. :7.9	Max. :4.4	Max. :6.9	Max. :2.5	

- The `plyr` package (Wickham, 2011) is very useful to manipulate data frames (df)
- Let's take a look at the `iris` built-in data frame

```
> summary(iris, digits = 2)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.3	Min. :2.0	Min. :1.0	Min. :0.1	setosa :50
1st Qu.:5.1	1st Qu.:2.8	1st Qu.:1.6	1st Qu.:0.3	versicolor:50
Median :5.8	Median :3.0	Median :4.3	Median :1.3	virginica :50
Mean :5.8	Mean :3.1	Mean :3.8	Mean :1.2	
3rd Qu.:6.4	3rd Qu.:3.3	3rd Qu.:5.1	3rd Qu.:1.8	
Max. :7.9	Max. :4.4	Max. :6.9	Max. :2.5	

- For instance, the `ddply` is applied on df and returns df, hence the dd
- Whereas, the `dply` is applied on df and returns a list, hence the dl

- The `plyr` package (Wickham, 2011) is very useful to manipulate data frames (`df`)
- Let's take a look at the `iris` built-in data frame

```
> summary(iris,digits = 2)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.3	Min. :2.0	Min. :1.0	Min. :0.1	setosa :50
1st Qu.:5.1	1st Qu.:2.8	1st Qu.:1.6	1st Qu.:0.3	versicolor:50
Median :5.8	Median :3.0	Median :4.3	Median :1.3	virginica :50
Mean :5.8	Mean :3.1	Mean :3.8	Mean :1.2	
3rd Qu.:6.4	3rd Qu.:3.3	3rd Qu.:5.1	3rd Qu.:1.8	
Max. :7.9	Max. :4.4	Max. :6.9	Max. :2.5	

- For instance, the `ddply` is applied on `df` and returns `df`, hence the `dd`
- Whereas, the `dply` is applied on `df` and returns a list, hence the `dl`
- To see this, we can look at the average sepal length for species using the following command

```
> library(plyr)
> ddply(iris,"Species", function(ds) mean(ds[, "Sepal.Length"]) )
```

	Species	V1
1	setosa	5.006
2	versicolor	5.936
3	virginica	6.588

- The `plyr` package (Wickham, 2011) is very useful to manipulate data frames (df)
- Let's take a look at the `iris` built-in data frame

```
> summary(iris,digits = 2)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.3	Min. :2.0	Min. :1.0	Min. :0.1	setosa :50
1st Qu.:5.1	1st Qu.:2.8	1st Qu.:1.6	1st Qu.:0.3	versicolor:50
Median :5.8	Median :3.0	Median :4.3	Median :1.3	virginica :50
Mean :5.8	Mean :3.1	Mean :3.8	Mean :1.2	
3rd Qu.:6.4	3rd Qu.:3.3	3rd Qu.:5.1	3rd Qu.:1.8	
Max. :7.9	Max. :4.4	Max. :6.9	Max. :2.5	

- For instance, the `ddply` is applied on df and returns df, hence the dd
- Whereas, the `dlply` is applied on df and returns a list, hence the dl
- To see this, we can look at the average sepal length for species using the following command

```
> library(plyr)
> ddply(iris,"Species", function(ds) mean(ds[, "Sepal.Length"]) )
```

	Species	V1
1	setosa	5.006
2	versicolor	5.936
3	virginica	6.588

- Otherwise, one can split the data based on species and work on each data separately

```
> iris.list <- dlply(iris,"Species",data.frame)
> length(iris.list)
```

```
[1] 3
```

```
> sapply(iris.list, nrow)
```

setosa	versicolor	virginica
50	50	50

## Practice

Using the `iris` dataset, which species has the highest maximum petal length-to-width ratio?

### Hints

- one can easily create a new variable called `XYZ` to dataset `ds` by assigning the corresponding values to a new called variable `ds[, "XYZ"] <- ....`
- use the `ddply` function

# Plots

- The main plot function in R is called `plot`
- The function is powerful and can be ran on different objects, for instance
  - use it on a dataset to get a perspective
  - use it on a regression for diagnostics
  - use it directly on time series objects, e.g. `xts`

# Plots

- The main plot function in R is called `plot`
- The function is powerful and can be ran on different objects, for instance
  - use it on a dataset to get a perspective
  - use it on a regression for diagnostics
  - use it directly on time series objects, e.g. `xts`
- For instance the plot can be ran on different objects, matrices, vectors, or data frames:

```
> y <- iris[,"Petal.Length"]  
> x <- iris[,"Petal.Width"]  
> plot(y~x)
```

- Alternatively, the same plot can be produced in one line command  

```
> plot(Petal.Length ~ Petal.Width, data = iris)
```

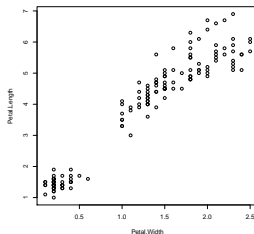


# Plots

- The main plot function in R is called `plot`
- The function is powerful and can be ran on different objects, for instance
  - use it on a dataset to get a perspective
  - use it on a regression for diagnostics
  - use it directly on time series objects, e.g. `xts`
- For instance the plot can be ran on different objects, matrices, vectors, or data frames:

```
> y <- iris[,"Petal.Length"]  
> x <- iris[,"Petal.Width"]  
> plot(y~x)
```
- Alternatively, the same plot can be produced in one line command

```
> plot(Petal.Length ~ Petal.Width,data = iris)
```
- In either case, we get



- In addition to `y` and `x`, one can specify a number of arguments
  - 1 `col`: setting color using an integer or character, e.g. `col = 2` or `col = "red"`
  - 2 `type`: type of plot. e.g. `type = "l"` returns a line
  - 3 `lwd`: width of line
  - 4 `lty`: type of line, e.g. dashed or solid
  - 5 `pch`: shape of dots
  - 6 `xlab`: x-axis labels
  - 7 `ylab`: y-axis labels
  - 8 `xlim`: range of x-axis (useful for multiple lines)
  - 9 `ylim`: range of y-axis (useful for multiple lines)

- In addition to `y` and `x`, one can specify a number of arguments
  - 1 `col`: setting color using an integer or character, e.g. `col = 2` or `col = "red"`
  - 2 `type`: type of plot. e.g. `type = "l"` returns a line
  - 3 `lwd`: width of line
  - 4 `lty`: type of line, e.g. dashed or solid
  - 5 `pch`: shape of dots
  - 6 `xlab`: x-axis labels
  - 7 `ylab`: y-axis labels
  - 8 `xlim`: range of x-axis (useful for multiple lines)
  - 9 `ylim`: range of y-axis (useful for multiple lines)
- After calling the plot, one can also add points, lines, or legend
  - 1 `lines`: adds line for a given `y` and `x`
  - 2 `points`: adds points for a given `y` and `x`
  - 3 `abline`: is useful to add horizontal, vertical, or linear lines

- In addition to `y` and `x`, one can specify a number of arguments

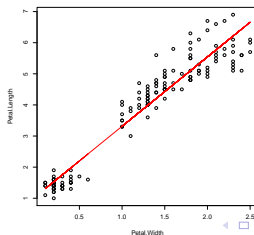
- 1 `col`: setting color using an integer or character, e.g. `col = 2` or `col = "red"`
- 2 `type`: type of plot. e.g. `type = "l"` returns a line
- 3 `lwd`: width of line
- 4 `lty`: type of line, e.g. dashed or solid
- 5 `pch`: shape of dots
- 6 `xlab`: x-axis labels
- 7 `ylab`: y-axis labels
- 8 `xlim`: range of x-axis (useful for multiple lines)
- 9 `ylim`: range of y-axis (useful for multiple lines)

- After calling the plot, one can also add points, lines, or legend

- 1 `lines`: adds line for a given `y` and `x`
- 2 `points`: adds points for a given `y` and `x`
- 3 `abline`: is useful to add horizontal, vertical, or linear lines

- For instance, one can add a fitted linear regression line as follows

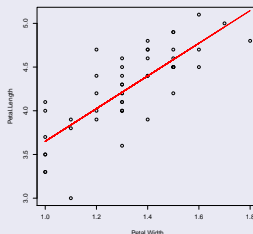
```
> iris[,"Petal.Length.hat"] <- fitted(lm(Petal.Length ~ Petal.Width,data = iris))  
> lines(Petal.Length.hat~Petal.Width, data = iris, col = 2)
```



## Practice

Write a function that plots the relationship between the petal length and width (in line with the figure from the previous slide), for a given species. The function should be named `lm.species` and should take a character as its main and only argument. For instance, the following command should yield

```
> lm.species("versicolor")
```



### Hints:

- 1 recall the `iris.list`
- 2 each item in `iris.list` corresponds to a sub-dataset
- 3 in lists, one need to use double brackets to retrieve a certain item
  - `iris.list[[1]]` returns the first dataset, which is the setosa species
  - alternatively, one can achieve so by `iris.list[["setosa"]]`

# Simulations

# RNG

- The basic sampling method in R is the `sample` command
  - allowing user to randomly sample a subset from a given set

```
> sample(1:10,5)
```

```
[1] 10 8 5 9 2
```

```
> sample(1:10,5)
```

```
[1] 9 1 3 6 4
```

# RNG

- The basic sampling method in R is the `sample` command
  - allowing user to randomly sample a subset from a given set

```
> sample(1:10,5)
```

```
[1] 10 8 5 9 2
```

```
> sample(1:10,5)
```

```
[1] 9 1 3 6 4
```

- One needs to specify whether the sampling is done with replacement

```
> sample(1:10,20)
```

```
Error in sample.int(length(x), size, replace, prob) :  
cannot take a sample larger than the population when 'replace = FALSE'
```

```
> sample(1:10,20,replace = T)
```

```
[1] 3 7 9 9 9 8 8 6 1 9 6 5 10 8 7 9 8 10 3 3
```



# RNG

- The basic sampling method in R is the `sample` command
  - allowing user to randomly sample a subset from a given set

```
> sample(1:10,5)
```

```
[1] 10 8 5 9 2
```

```
> sample(1:10,5)
```

```
[1] 9 1 3 6 4
```

- One needs to specify whether the sampling is done with replacement

```
> sample(1:10,20)
```

```
Error in sample.int(length(x), size, replace, prob) :  
cannot take a sample larger than the population when 'replace = FALSE'
```

```
> sample(1:10,20,replace = T)
```

```
[1] 3 7 9 9 9 8 8 6 1 9 6 5 10 8 7 9 8 10 3 3
```

- To set a seed, one can use `set.seed` function

```
> apply(1:5, function(i) { set.seed(13); sample(1:10,5) })
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	8	8	8	8	8
[2,]	3	3	3	3	3
[3,]	4	4	4	4	4
[4,]	1	1	1	1	1
[5,]	6	6	6	6	6

# Distributions

- Most distributions in R follow similar commands
- Assume we are interested in the normal distribution, then
  - 1 `dnorm` is the probability density function
  - 2 `pnorm` is the cumulative distribution function
  - 3 `qnorm` is the quantile for a given normal
  - 4 `rnorm` is a RNG for a given `n` sample

# Distributions

- Most distributions in R follow similar commands
- Assume we are interested in the normal distribution, then
  - 1 `dnorm` is the probability density function
  - 2 `pnorm` is the cumulative distribution function
  - 3 `qnorm` is the quantile for a given normal
  - 4 `rnorm` is a RNG for a given `n` sample
- In fact, our focus will be mainly on the RNG functions for normal distribution
- By default, the `rnorm` is a standard normal, i.e.

```
> rnorm
```

```
function (n, mean = 0, sd = 1)
```

# Distributions

- Most distributions in R follow similar commands
- Assume we are interested in the normal distribution, then
  - 1 `dnorm` is the probability density function
  - 2 `pnorm` is the cumulative distribution function
  - 3 `qnorm` is the quantile for a given normal
  - 4 `rnorm` is a RNG for a given `n` sample
- In fact, our focus will be mainly on the RNG functions for normal distribution
- By default, the `rnorm` is a standard normal, i.e.

```
> rnorm  
  
function (n, mean = 0, sd = 1)
```

- Let  $X_i \sim N(i, 1)$  iid, for  $i = 1, \dots, 5$ ,
- To simulate  $n = 10^3$  observations for each variable, one can run

```
> norm.list <- lapply(1:5, function(i) rnorm(10^3, mean = i, sd = 1) )  
> sapply(norm.list, mean)
```

```
[1] 1.039613 2.012433 3.068028 4.015763 4.812596
```

## Practice

- Let  $X_1$  and  $X_2$  be two independent random Gaussian variables, with  $\mu_1 = 10$ ,  $\mu_2 = 15$ ,  $\sigma_1 = 3$ , and  $\sigma_2 = 5$ , i.e.

$$X_1 \sim N(10, 3^2) \quad (1)$$

and

$$X_2 \sim N(15, 5^2) \quad (2)$$

- Using a simulation, find the probability that

$$\mathbb{P}(X_1 < X_2) \quad (3)$$

- In addition, compare your result with the exact solution.

- Hint:** since  $X_1$  and  $X_2$  are independent, then

$$X_1 - X_2 \sim N(\mu_1 - \mu_2, \sigma_1^2 + \sigma_2^2) \quad (4)$$

# Multivariate Normal

- We know when  $X_i \sim N(\mu_i, \sigma_i^2)$ , it follows that

$$X = \begin{bmatrix} X_1 \\ \vdots \\ X_d \end{bmatrix} \sim N_d(\mu, \Sigma) \quad (5)$$

with

- $\underline{\mu}$  denoting a  $d \times 1$  vector of mean returns and
- $\Sigma$  is a non-singular  $d \times d$  covariance matrix

# Multivariate Normal

- We know when  $X_i \sim N(\mu_i, \sigma_i^2)$ , it follows that

$$X = \begin{bmatrix} X_1 \\ \vdots \\ X_d \end{bmatrix} \sim N_d(\mu, \Sigma) \quad (5)$$

with

- $\mu$  denoting a  $d \times 1$  vector of mean returns and
- $\Sigma$  is a non-singular  $d \times d$  covariance matrix
- One can simulate the vector  $X$  using the `mvrnorm` function from the MASS package (Venables & Ripley, 2002)
- Nevertheless, one needs to specify a vector of means,  $\mu$ , and a covariance matrix,  $\Sigma$

# Multivariate Normal

- We know when  $X_i \sim N(\mu_i, \sigma_i^2)$ , it follows that

$$X = \begin{bmatrix} X_1 \\ \vdots \\ X_d \end{bmatrix} \sim N_d(\mu, \Sigma) \quad (5)$$

with

- $\mu$  denoting a  $d \times 1$  vector of mean returns and
- $\Sigma$  is a non-singular  $d \times d$  covariance matrix
- One can simulate the vector  $X$  using the `mvrnorm` function from the MASS package (Venables & Ripley, 2002)
- Nevertheless, one needs to specify a vector of means,  $\mu$ , and a covariance matrix,  $\Sigma$

- Let's get back to the EuStockMarkets dataset

```
> ds <- EuStockMarkets
> R <- (ds[-1,]/ds[-nrow(ds),] - 1)*100 # computes returns
> Mu <- apply(R, 2, mean) # mean vector
> Sigma <- var(R) # covariance matrix
```

- For a given  $\mu$  and  $\Sigma$ , one can generate  $n$  vectors of  $X$

```
> n <- 10^3
> X <- mvrnorm(N, Mu, Sigma)
> dim(X)
```

```
[1] 1000    4
```

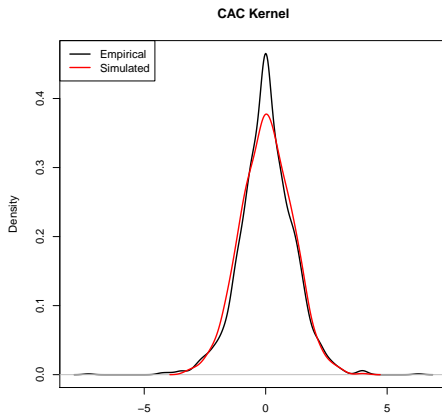


- Let's compare between the simulated distribution and the empirical one
- The density function allows the user to estimate the kernel density of variable  $x$

- Let's compare between the simulated distribution and the empirical one
- The density function allows the user to estimate the kernel density of variable  $x$

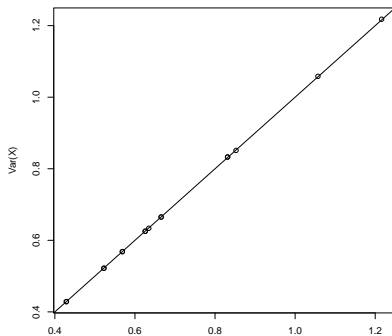
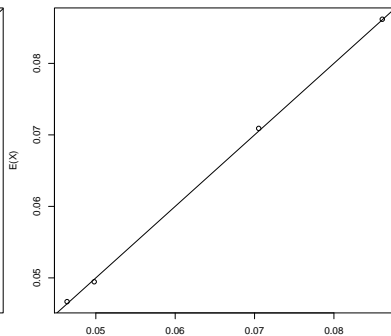
- Looking at the CAC index, we have

```
> kernel_emp <- density(R["CAC"])  
> kernel_sim <- density(X["CAC"])  
> plot(kernel_emp, main = "CAC Kernel", lwd = 2)  
> lines(kernel_sim, col = 2, lwd = 2)  
> legend("topleft", c("Empirical", "Simulated"), col = 1:2, lwd = 2)
```



- Also, note since  $X$  is generated using  $\mu$  and  $\Sigma$ , the mean and covariance of the simulated data should be consistent

```
> n <- 10^6  
> X <- mvrnorm(n, Mu, Sigma)  
> plot(var(X) ~ Sigma, xlab = expression(Sigma), ylab = "Var(X)")  
> abline(a = 0, b = 1)  
> plot(apply(X, 2, mean) ~ Mu, xlab = expression(mu), ylab = "E(X)")  
> abline(a = 0, b = 1)
```

(a)  $\Sigma$ (b)  $\mu$

## Practice

- Consider a portfolio strategy that equally allocates among the 4 indices from EuStock-Markets
  - assume daily re-balancing
- Simulate the portfolio return over 250 days
- Repeat the above step 1000 times and draw the density of the portfolio Sharpe-ratio (SR)
- Make some observations about the performance

## Practice

- Consider a portfolio strategy that equally allocates among the 4 indices from EuStock-Markets
  - assume daily re-balancing
- Simulate the portfolio return over 250 days
- Repeat the above step 1000 times and draw the density of the portfolio Sharpe-ratio (SR)
- Make some observations about the performance

## Hint

Create a function `SR_f` that simulates the portfolio return over 250 days and returns. Given this, compute the SR using the average to standard deviation ratio.

```
> SR_f <- function(i) {  
+   R.250 <- mvrnorm(250,Mu,Sigma)  
+   Rp <- R.250*%W  
+   SR <- (mean(Rp)/sd(Rp))  
+   return(SR)  
+ }
```

## Practice

- Consider a portfolio strategy that equally allocates among the 4 indices from EuStock-Markets
  - assume daily re-balancing
- Simulate the portfolio return over 250 days
- Repeat the above step 1000 times and draw the density of the portfolio Sharpe-ratio (SR)
- Make some observations about the performance

## Hint

Create a function `SR_f` that simulates the portfolio return over 250 days and returns. Given this, compute the SR using the average to standard deviation ratio.

```
> SR_f <- function(i) {
+   R.250 <- mvrnorm(250,Mu,Sigma)
+   Rp <- R.250*%W
+   SR <- (mean(Rp)/sd(Rp))
+   return(SR)
+ }
```

## Note

If  $W$  denotes a  $d \times 1$  vector of portfolio weights while  $R$  is a  $n \times d$  denoting the returns of  $d$  assets over  $n$  periods, then the portfolio return is given by a  $n \times 1$  vector  $R_p$ , such that

$$R_p = RW \quad (6)$$

For matrix multiplication, use the `%*` operator

# Cholesky Decomposition

- For the sake of argument, assume we cannot use the `mvnrm` function
- How can we, then, simulate a multivariate Gaussian data?
- One answer is the Cholesky Decomposition (CD)

- The idea behind CD is the following

- 1 Simulate  $d$  univariate standard normally distributed random variables and stack it in a  $n \times d$  matrix called  $Z$
- 2 For a given  $\Sigma$ , use the CD to find matrix  $A$ , such that

$$\Sigma = AA' \quad (7)$$

- 3 Finally, map the  $Z$  matrix with respect to  $A$  and  $\mu$

$$X = M + ZA \quad (8)$$

where  $M$  is a  $n \times d$  matrix with each row is equal to  $\mu'$

## CD in Action

- Let's consider again the EuStockMarkets data
- Since we have 4 indices, then  $d = 4$
- Simulate 4 standard normal random variables over  $n = 10^6$  sample, named  $Z$
- Since we know  $\Sigma$ , use `chol` function to find  $A$
- Finally, map  $Z$  using  $\mu$  and  $A$

```
> n <- 10^6  
> A <- chol(Sigma)  
> M <- matrix(Mu,n,length(Mu),byrow = T)  
> Z <- sapply( 1:length(Mu),function(i) rnorm(n) )  
> X_CD <- M + Z%*%A
```



## CD in Action

- Let's consider again the EuStockMarkets data
- Since we have 4 indices, then  $d = 4$
- Simulate 4 standard normal random variables over  $n = 10^6$  sample, named  $Z$
- Since we know  $\Sigma$ , use `chol` function to find  $A$
- Finally, map  $Z$  using  $\mu$  and  $A$

```
> n <- 10^6  
> A <- chol(Sigma)  
> M <- matrix(Mu, n, length(Mu), byrow = T)  
> Z <- sapply( 1:length(Mu), function(i) rnorm(n) )  
> X_CD <- M + Z%*%A
```

## Reflect

- Compare between  $X\_CD$  and  $X$
- Recall that  $X$  was produced using the `mvrnorm` function
- Both should yield similar results
  - in terms of moments and kernel

# Option Pricing

# Simulating Stock Prices

- If the stock price obeys to the following dynamics

$$\frac{dS_t}{S_t} = \mu t + \sigma dW_t \quad (9)$$

then for an initial price  $S_0$ , the solution is given by

$$\log\left(\frac{S_t}{S_0}\right) = \left(\mu - \frac{\sigma^2}{2}\right)t + \sigma W_t \quad (10)$$

with

$$W_t \sim N(0, t) \quad (11)$$

# Simulating Stock Prices

- If the stock price obeys to the following dynamics

$$\frac{dS_t}{S_t} = \mu t + \sigma dW_t \quad (9)$$

then for an initial price  $S_0$ , the solution is given by

$$\log\left(\frac{S_t}{S_0}\right) = \left(\mu - \frac{\sigma^2}{2}\right)t + \sigma W_t \quad (10)$$

with

$$W_t \sim N(0, t) \quad (11)$$

- The log of the ratio  $S_t/S_0$  denotes the change in price over time  $t$ , i.e. return

# Simulating Stock Prices

- If the stock price obeys to the following dynamics

$$\frac{dS_t}{S_t} = \mu dt + \sigma dW_t \quad (9)$$

then for an initial price  $S_0$ , the solution is given by

$$\log\left(\frac{S_t}{S_0}\right) = \left(\mu - \frac{\sigma^2}{2}\right)t + \sigma W_t \quad (10)$$

with

$$W_t \sim N(0, t) \quad (11)$$

- The log of the ratio  $S_t/S_0$  denotes the change in price over time  $t$ , i.e. return
- Let  $r_t$  denote the return on the stock price over  $t$ , such that

$$r_t \sim N\left(t\left(\mu - \frac{\sigma^2}{2}\right), t\sigma^2\right) \quad (12)$$

# Simulating Stock Prices

- If the stock price obeys to the following dynamics

$$\frac{dS_t}{S_t} = \mu dt + \sigma dW_t \quad (9)$$

then for an initial price  $S_0$ , the solution is given by

$$\log\left(\frac{S_t}{S_0}\right) = \left(\mu - \frac{\sigma^2}{2}\right)t + \sigma W_t \quad (10)$$

with

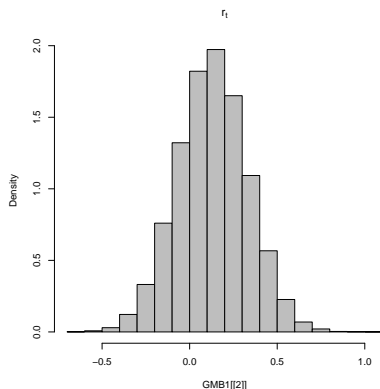
$$W_t \sim N(0, t) \quad (11)$$

- The log of the ratio  $S_t/S_0$  denotes the change in price over time  $t$ , i.e. return
- Let  $r_t$  denote the return on the stock price over  $t$ , such that

$$r_t \sim N\left(t\left(\mu - \frac{\sigma^2}{2}\right), t\sigma^2\right) \quad (12)$$

- Since we worked with normal distributions, we know how to simulate  $W_t$
- From Equation (10), if we know how to simulate  $W_t$ , then we can simulate either  $r_t$  or  $S_t$

- ```
> GMB_f <- function(n,S0,Mu,Sigma,Time) {
+   EX <- Time*(Mu - (Sigma^2)/2)
+   VX <- Time*(Sigma^2)
+   r_t <- rnorm(n,EX,sqrt(VX))
+   S_t <- S0*exp(r_t)
+   list(S_t,r_t)
+ }
> GMB1 <- GMB_f(10^5,100,0.15,0.2,1)
> hist(GMB1[[1]], main = expression(S[t]), col = "gray", freq = F)
> hist(GMB1[[2]], main = expression(r[t]), col = "gray", freq = F)
```



# Simulating Multiple Stock Prices

- Simulating multiple stock prices follow suit with the procedure from (5)
- Equation (12) takes the following multivariate form

$$r_t = \begin{bmatrix} r_{t,1} \\ \vdots \\ r_{t,d} \end{bmatrix} \sim N_d(t\mu, t\Sigma) \quad (13)$$

- Note that the  $i$ th of the mean vector  $\mu_t$  is given by

$$\mu_i = \left( \mu_i - \frac{\sigma_i^2}{2} \right), \forall i \in \{1, \dots, d\} \quad (14)$$

and the  $i$ th row and  $j$ th column of  $\Sigma_t$  is

$$\Sigma_{t,ij} = \rho_{ij}\sigma_i\sigma_j \quad (15)$$

where  $\rho_{ij}$  is the correlation coefficient between  $i$  and  $j$ ,  $\forall i, j \in \{1, \dots, d\}$  and  $\rho_{ij} = 1$   $\forall i = j$ .

**Ultimately**, if we know  $\mu$  and  $\Sigma$ , we can simulate terminal prices for  $d$  stocks



- Let's consider the following case with three stocks
  - $S_{0,1} = S_{0,2} = S_{0,3} = 100$
  - $\sigma_1 = 0.2$ ,  $\sigma_2 = 0.3$ , and  $\sigma_3 = 0.25$
  - For  $i = 1, 2, 3$ , we have  $\mu_i = 0.06$
  - $\rho_{12} = 0.5$ ,  $\rho_{13} = 0.25$ , and  $\rho_{23} = -0.25$

### Practice

- Simulate the prices of the three stocks for  $t = 1$  year
- Price the following spread option between stock 1 and 2 (i.e. Q9 from HW2)

- Let's consider the following case with three stocks
  - $S_{0,1} = S_{0,2} = S_{0,3} = 100$
  - $\sigma_1 = 0.2$ ,  $\sigma_2 = 0.3$ , and  $\sigma_3 = 0.25$
  - For  $i = 1, 2, 3$ , we have  $\mu_i = 0.06$
  - $\rho_{12} = 0.5$ ,  $\rho_{13} = 0.25$ , and  $\rho_{23} = -0.25$

### Practice

- Simulate the prices of the three stocks for  $t = 1$  year
- Price the following spread option between stock 1 and 2 (i.e. Q9 from HW2)

### Hint

Recall that the covariance matrix can be decomposed as

$$\Sigma = \mathbf{D}\mathbf{R}\mathbf{D} \quad (16)$$

where

- $\mathbf{D}$  is the diagonal matrix of standard deviations, i.e.  $\mathbf{D}_{ii} = \sigma_i \forall i = 1, \dots, d$  and  $\mathbf{D}_{ij} = 0 \forall i \neq j$
- $\mathbf{R}$  is the correlation matrix, where  $\mathbf{R}_{ij} = \rho_{ij}$  and  $\mathbf{R}_{ii} = 1 \forall i = 1, \dots, d$

- The main issue is identifying the parameters from (13)
- If we can, then simulating the prices is straightforward
  - with or without the `mvnrm` function

- The main issue is identifying the parameters from (13)
- If we can, then simulating the prices is straightforward
  - with or without the `mvrnorm` function

```
> Time <- 1
> r <- 0.06
> K <- 1
> Mu <- rep(r,3)
> D <- diag(c(0.2,0.3,0.5)) # diagonal of sigmas
> R <- rbind(c(1,0.5,0.25), c(0.5,1,-0.25), c(0.25,-0.25,1))
> Sigma <- D%*%R%*%D
> S <- rep(100,3)
> n <- 10^4
> GMB_MV_f <- function(n,S,Mu,Sigma,Time) {
+   EX <- Time*(Mu - (diag(D)^2)/2)
+   VX <- Time*(Sigma)
+   r_t <- mvrnorm(n,EX,VX)
+   S_t <- sapply(1:ncol(r_t), function(i) exp(r_t[,i])*S[i])
+   list(S_t,r_t)
+ }
> GMB2 <- GMB_MV_f(n,S,Mu,Sigma,Time)
> S_T <- GMB2[[1]]
> Spread <- S_T[,1] - S_T[,2] - K
> Spread[Spread < 0] <- 0
> mean(Spread)*exp(-r*Time)
```

```
[1] 10.06153
```

# Open Source Data

# References I

- [ ]Ang, C. S. 2015. *Analyzing financial data and implementing financial models using r*. Springer.
- [ ]James, G., Witten, D., Hastie, T., & Tibshirani, R. 2014. An introduction to statistical learning: with applications in r.
- [ ]Venables, W. N., & Ripley, B. D. 2002. *Modern applied statistics with s* (Fourth ed.). New York: Springer. Retrieved from <http://www.stats.ox.ac.uk/pub/MASS4> (ISBN 0-387-95457-0)
- [ ]Wickham, H. 2011. The split-apply-combine strategy for data analysis. *Journal of Statistical Software*, 40(1), 1–29. Retrieved from <http://www.jstatsoft.org/v40/i01/>